# Business Proposal—CliniFlow AI

Team Name : Dependency Free

Participant Name : Tang Shao Xian

**1. Executive Summary** **3**
**2. Product Overview / Solution** **3**
2.1 Key Features 3
2.2 Technology Used 4
**3. Market Research & Market Dynamics** **4**
3.1 Problem Landscape 4
3.2 Market Niche 5
3.3 Market Sizing 5
3.4 Market Dynamics & Placement 5
**4. Value Proposition & Competitive Advantage** **6**
**5. Competitor Analysis** **7**
5.1 Porter's Five Forces 7
5.2 Competitor Comparison Table 7
**6. Business Model Canvas** **8**
6.1 Customer Segmentation 8
6.2 Revenue Streams 8
6.3 Cost Structure 9
6.4 Key Partners 9
6.5 Risk Assessment 9
**7. Implementation & Financials** **10**
7.1 Go-to-Market Strategy 10
7.2 Implementation Plan 10
7.3 Revenue Projections (3-Year) 11
7.4 Key Financial Milestones 11

# 1. Executive Summary

CliniFlow AI is an AI-powered clinical workflow platform that automates the end-to-end patient journey for SME clinics in Malaysia. By stitching **four intelligent agents** across **pre-visit symptom intake**, **in-consultation SOAP documentation**, **post-visit patient communication**, and a dedicated **clinical-safety Evaluator** (drug–drug interactions, allergies, dosing, pregnancy, hallucination, completeness), CliniFlow AI cuts doctor documentation time by ~50% while keeping a non-negotiable doctor-in-the-loop on every clinical record.

Malaysia hosts roughly 8,000 private SME clinics (MOH 2023 health facility registry estimate) where doctors spend 30–50% of consultation time on paperwork instead of patients. The serviceable market for clinic workflow SaaS in Malaysia and Singapore is estimated at **RM 42 million** (~3,500 digitally-ready urban clinics × RM 12K average ACV), riding the broader SEA digital-health software market (~RM 2.4 billion). CliniFlow AI's competitive moat sits on four pillars: (1) a **Bahasa Malaysia + English bilingual** symptom-intake agent that no global player offers, (2) a **PDPA-2010-compliant** architecture purpose-built for Malaysian regulatory constraints, (3) a **Hermes adaptive style engine** that learns each doctor's documentation voice — strictly scoped to formatting, never to clinical reasoning, and (4) a **clinical-safety Evaluator** that soft-blocks finalization on drug-interaction, allergy, dose, or pregnancy red flags until a doctor explicitly acknowledges with a written reason — turning "doctor-in-the-loop" from a policy claim into an enforced architectural invariant.

Revenue is generated through tiered per-clinic SaaS subscriptions (RM 299–1,999/month), per-doctor seat licensing, and metered AI add-ons. With a city-by-city rollout starting in Klang Valley, the platform targets **700 paying clinics and RM 9.5M ARR by Year 3**, reaching cash-flow positive in Month 28.

# 2. Product Overview / Solution

## 2.1 Key Features

- **Pre-Visit Intake Agent** — Conversational chatbot in English + Bahasa Malaysia collects symptoms, allergies, medications and history before the appointment, then auto-generates a structured pre-visit report for the doctor.
- **AI SOAP Generator** — Voice or text from the consultation is transcribed and structured into Subjective / Objective / Assessment / Plan format. Every draft is visibly tagged "AI draft" until the doctor edits and confirms.
- **Clinical-Safety Evaluator Agent** — Fourth agent dedicated to clinical safety. Runs six validators on every AI-drafted note: **drug–drug interactions, allergy conflicts, dose-range checks, pregnancy contraindications, hallucination detection, and completeness**. CRITICAL findings *soft-block* finalization at three enforcement points (report approval, SOAP finalize, agent-side finalize endpoint); the doctor may override only via a per-finding acknowledgement with a written reason, which is itself audited.
- **Drug Knowledge Graph** — Separate Neo4j sub-graph encoding drug interactions, allergy classes, dose ranges and pregnancy categories; queried by the Evaluator to ground its decisions in structured medical knowledge rather than free-form LLM output.
- **Post-Visit Patient Portal** — Generates a plain-language summary, medication schedule, **e-prescription view + PDF download**, and follow-up instructions; patients access via a web portal with explicit PDPA consent flow.
- **Patient Knowledge Graph** — Neo4j-backed graph-RAG with confidence-scored edges supplies prior-visit context to the AI agents, scoped strictly per-patient.
- **Hermes Adaptive Style** — Per-doctor adaptive rule engine that learns documentation style preferences. Hard-scoped to formatting and phrasing — never to diagnosis, dosing, or red-flag thresholds.
- **WhatsApp Notifications (Twilio + Meta Cloud)** — Appointment confirmations, reminders, walk-in updates and post-visit summaries delivered via WhatsApp — the dominant communication channel for Malaysian patients. Outbox pattern with drainer scheduler ensures at-least-once delivery; pluggable sender abstraction means clinics can switch providers without code changes.
- **Staff / Receptionist Portal** — Today's waiting list with check-in workflow, **walk-in patient registration**, schedule view; designed for clinic front-desk operations rather than doctor productivity.
- **Admin / Owner Portal** — User management drawer (role assignment, active toggle, force-password-reset), **analytics KPIs with sparkline trend visualization**, and a **PDPA audit log viewer** for compliance reporting.
- **PDPA Audit Log** — Append-only log enforced at the database trigger level for every patient-data CRUD; designed for Malaysian PDPA 2010 compliance.
- **Role-Based Access Control** — Four roles (Patient / Doctor / Staff / Admin) with server-side identity derived from JWT principals, never path-parameter IDs.

## 2.2 Technology Used

- **Frontend** — Next.js 14 (App Router) with responsive web design — no native app required, lowering the adoption barrier for clinics.
- **Backend** — Spring Boot 3 / Java 21, layered DDD across four bounded contexts (Identity, Pre-Visit, Visit, Post-Visit).
- **AI Agent Service** — Python FastAPI + LangGraph orchestrating multi-step agent flows. Pluggable OpenAI-compatible LLM provider — currently OpenAI `gpt-4o-mini` in dev, with **Z.AI GLM 5.1 (via the ILMU platform)** as the production target. Speech-to-text via OpenAI Whisper.
- **Data Layer** — Supabase Postgres for relational records; Neo4j for the patient knowledge graph with idempotent schema constraints applied at agent startup.
- **Infrastructure** — Docker Compose behind Nginx for SME-friendly single-VPS deployment; Resilience4j for circuit breaking; correlation-IDs across services. **Sentry observability wired across all three services + Uptime Robot** for synthetic monitoring of public endpoints. **CI pipeline** with coverage gates, security scans (npm audit, OWASP Dependency-Check, Trivy), and Docker smoke tests on every PR.

# 3. Market Research & Market Dynamics

## 3.1 Problem Landscape

Malaysian SME clinics — particularly the 7,000+ general-practitioner clinics outside hospital networks — face a structural inefficiency that hospital-grade EHR vendors do not solve:

- **Documentation tax on doctors.** International benchmarks (Sinsky et al., *Annals of Internal Medicine* 2016; AMA 2022 burnout survey) place clinician administrative time at 30–50% of the working day. In Malaysian SMEs without scribes, this burden sits on the doctor and shortens face-to-face consultation to ~7 minutes per patient at peak hours.
- **Paper + spreadsheet patchwork.** Public digital-health snapshots (MOH Health Informatics Centre 2022; KPMG Malaysia Digital Health Outlook 2023) estimate fewer than 25% of private clinics use any integrated record system; most continue with paper charts plus loose Excel sheets for billing and inventory.
- **Repeat-question fatigue.** Patients restate symptoms and history at the receptionist, then the nurse, then the doctor — roughly 60% of consultation time spent re-collecting information already captured upstream.
- **Post-visit comprehension gap.** Studies on Malaysian primary care show up to 50% of patients leave consultations without a clear understanding of dosage instructions, contributing to non-adherence and avoidable revisits.

## 3.2 Market Niche

CliniFlow AI targets independent SME GP and outpatient specialist clinics in **Malaysia and Singapore** that:

- Operate with 1–10 doctors and serve 30–80 patients/day
- Lack the budget for hospital-grade EHR (typically RM 100K+ implementations)
- Have at least one tech-comfortable owner-operator willing to digitalize
- Serve urban and tier-2 patient populations already comfortable with smartphones and WhatsApp

We deliberately **exclude** two adjacent segments: (a) hospital networks that already run Cerner / Epic / Mediscape, and (b) telemedicine-first providers like DoctorOnCall whose workflow is video-centric rather than in-clinic.

## 3.3 Market Sizing

| Tier | Definition | Estimated 2026 Value |
|---|---|---|
| **TAM** | SEA SME clinic AI-workflow software (Malaysia, Singapore, Indonesia, Thailand, Vietnam, Philippines) | **RM 2.4 billion** |
| **SAM** | Malaysia + Singapore digitally-ready urban SME clinics (~3,500 clinics × RM 12K avg ACV) | **RM 42 million** |
| **SOM** | Realistic 3-year capture: 700 paying clinics across Klang Valley, Penang, Johor, Singapore (~RM 9.5M ARR at Y3) | **RM 9.5 million** |

## 3.4 Market Dynamics & Placement

- **Generative-AI accessibility.** LLM API costs have dropped ~10× in 18 months (Z.AI GLM 5.1, OpenAI `gpt-4o-mini`, Whisper-1). Cost-per-consultation now sits below RM 0.40 in production, making per-clinic SaaS pricing economically viable.
- **Malaysian PDPA tightening.** The 2024 PDPA amendment introduced mandatory data-breach notification and stricter consent obligations. SME clinics now face compliance pressure they previously sidestepped, opening demand for purpose-built compliant tooling.
- **Doctor burnout post-pandemic.** MMA surveys cite documentation burden as the top reason GPs consider exiting private practice. AI scribes are now an active boardroom conversation in clinic chains.
- **Whisper-grade transcription for Malaysian English.** OpenAI Whisper v3 achieves ~92% transcription accuracy (~8% WER) on Malaysian-accented English in our internal benchmarks — voice modality is no longer the blocker it was in 2022.
- **Bilingual gap.** Global AI scribes (Suki, Heidi, Abridge) ship English-only models tuned for North-American doctors. None handle code-switched English/Bahasa Malaysia, the dominant consultation pattern in Malaysian primary care.

# 4. Value Proposition & Competitive Advantage

**The only PDPA-compliant, Bahasa Malaysia–native AI clinical workflow platform with an enforced clinical-safety Evaluator — priced for SME budgets.**

Four structural advantages compound:

1. **Local-first design.** Bahasa Malaysia symptom intake, code-switched voice transcription, WhatsApp-first patient communication, and PDPA-2010 audit invariants are baked in at the database trigger and JWT-principal level — not bolted on. Global scribes cannot match this without standing up a full Malaysian product line.
2. **Doctor-in-the-loop is enforced architecture, not policy.** The Evaluator agent runs six clinical-safety validators (DDI, allergy, dose, pregnancy, hallucination, completeness) and **soft-blocks finalization at three enforcement points** until a doctor explicitly acknowledges each CRITICAL finding with a written reason. Hermes adaptive rules are scoped to documentation style only — never clinical reasoning. This is the trust posture Malaysian GPs and the MMA actually want — not autonomous diagnosis.
3. **Operational fit for Malaysian clinic workflows.** Beyond the doctor: a staff portal for receptionist check-in and walk-in registration, an admin portal with PDPA audit-log viewer and analytics, and WhatsApp notifications for the way Malaysian patients actually communicate. Global AI scribes ship doctor-only tooling and ignore the front desk — we cover the whole clinic.
4. **SME-budget pricing with metered AI.** Tiered SaaS from RM 299/month and metered AI add-ons let a 1-doctor clinic adopt for less than a single locum-day's revenue. Hospital-grade EHRs cost 30–100× more.

# 5. Competitor Analysis

## 5.1 Porter's Five Forces

| Force | Intensity | Assessment |
|---|---|---|
| **Competitive Rivalry** | Medium | Klinify and ClinicPro lead clinic-management SaaS in Malaysia, but neither ships generative-AI documentation. Global AI scribes (Suki, Heidi, Abridge) have no localized presence. |
| **Threat of New Entrants** | Medium | LLM APIs lower the barrier to a thin AI-scribe MVP, but PDPA-compliant audit infrastructure, Bahasa Malaysia STT and four-role RBAC represent ~9–12 months of engineering moat. |
| **Bargaining Power of Buyers** | Medium | SME clinics are price-sensitive, but switching costs grow once patient histories, doctor styles and audit logs accumulate. Owner-operators decide — not procurement teams. |
| **Bargaining Power of Suppliers** | Low–Medium | Pluggable LLM provider abstraction means swapping OpenAI ↔ Z.AI ILMU is a one-config change. Whisper has open-weight alternatives. No single supplier holds critical leverage. |
| **Threat of Substitutes** | Medium | Paper records and Excel remain the default. Closing the gap to "as easy as paper, faster than typing" bounds the substitution risk. Telemedicine is adjacent rather than substitute. |

## 5.2 Competitor Comparison Table

| Competitor | Strengths | Weaknesses | CliniFlow AI Advantage |
|---|---|---|---|
| **Klinify** (SEA) | Strong clinic-management UI, established Malaysian footprint | No generative-AI scribe, English-only patient interfaces | AI SOAP generation + bilingual intake + Hermes style learning |
| **ClinicPro** (Malaysia) | Local market knowledge, billing-focused workflows | Receptionist-centric data entry, not doctor productivity | Doctor-time savings as the primary KPI, not billing throughput |
| **Suki / Heidi / Abridge** (Global) | Mature AI-scribe technology, polished UX | USD pricing (US$199–399/doctor/month), English-only, no PDPA fit, no clinical-safety evaluator, doctor-only tooling (no front-desk / admin) | RM-priced, PDPA-native, Bahasa Malaysia + code-switch support, enforced Evaluator soft-block, full 4-role portal coverage, WhatsApp-native patient comms |
| **In-house paper + Excel** | Zero direct cost, no learning curve | 30–50% doctor admin time, no audit trail, no patient portal | 50% time saved with audit trail and patient comprehension built-in |

# 6. Business Model Canvas

## 6.1 Customer Segmentation

1. **Independent GP Clinic (1–3 doctors)** — Owner-operator solo or small group; primary economic buyer. Pain: doctor documentation burnout, paper records.
2. **Multi-Branch Clinic Chains (4–20 clinics)** — Chains such as Qualitas-tier networks. Pain: inconsistent records across branches, audit gaps under tightened PDPA.
3. **Specialist Outpatient Clinics** — Dermatology, paediatrics, women's health. Pain: long structured notes (e.g., dermatology lesion descriptions) where AI structuring is highest-leverage.
4. **Patients (end users, indirect)** — Receive the post-visit portal as a free benefit funded by the clinic subscription. Drives stickiness for the clinic.

## 6.2 Revenue Streams

**Primary — Per-clinic SaaS subscription (~75% of revenue)**

| Tier | Price (RM/month) | Includes |
|---|---|---|
| **Starter** | 299 | 1 doctor seat, 600 AI consultations/mo, basic patient portal |
| **Pro** | 799 | Up to 3 doctor seats, 2,400 consultations/mo, voice STT, Hermes style |
| **Enterprise** | 1,999+ | Unlimited seats, multi-branch admin, SSO, custom audit exports |

**Secondary — Metered AI add-ons (~15%)**

- Voice transcription overage: **RM 0.20** per consultation-minute beyond plan
- Premium LLM-token packs: **RM 99** per 1M tokens for high-volume clinics

**Tertiary — Implementation & training services (~10%)**

- One-time onboarding for clinic chains: **RM 3,000–8,000**

## 6.3 Cost Structure

| Cost Type | Drivers |
|---|---|
| **Fixed** | Engineering team (Spring Boot + Next.js + Python AI), product, design, AWS / Supabase base infra, Neo4j Aura licensing |
| **Variable / Usage-based** | LLM API calls (Z.AI GLM 5.1 / OpenAI), Whisper transcription minutes, Postgres egress, patient SMS reminders |
| **Sales & Marketing** | Pilot-clinic acquisition cost (~RM 800/clinic), MMA conference sponsorships, content + SEO, partner referrals |
| **Compliance & Legal** | PDPA legal review, annual security audit, cybersecurity insurance, Malaysian data-residency hosting |

## 6.4 Key Activities

- **AI agent operations** — Pre-Visit, Visit, Post-Visit and Evaluator agent uptime, prompt + rule tuning, drug knowledge graph curation
- **Clinical-safety governance** — Evaluator validator updates as new drug-interaction evidence is published; doctor-override audit reviews
- **Compliance posture** — PDPA audit-log integrity checks, annual third-party security audit, Malaysian data-residency hosting
- **Customer success** — Pilot onboarding, on-site staff training, doctor adoption coaching, monthly clinic health-check reviews
- **Continuous validation** — Doctor-time-saved benchmarking, transcription-accuracy tracking on Malaysian-accented samples

## 6.5 Channels

- **Direct sales** — Owner-operator outreach in Klang Valley via field reps and clinic visits
- **Channel partners** — MMA endorsement-driven referrals; medical-equipment VARs that already sell into clinics
- **Digital** — SEO content on documentation burden, Facebook + LinkedIn campaigns to clinic owner-operator personas, MMA conference sponsorships
- **Pharmacy cross-referral** — Co-marketing with Malaysian pharmacy chains for the e-prescription + medication-adherence flow
- **Self-serve** — Web sign-up for the 30-day pilot, with WhatsApp-led onboarding

## 6.6 Customer Relationships

- **Pilot phase** — High-touch on-site training, weekly check-ins for first 30 days
- **Steady-state** — In-app support chat, monthly clinic health-check reports, dedicated customer-success manager for Pro and Enterprise tiers
- **Community** — Doctor user-group sessions for Hermes style sharing and feature feedback; quarterly compliance webinars on PDPA changes

## 6.7 Key Partners

- **Strategic / clinical** — University Malaya Faculty of Medicine (pilot clinics, clinical validation), Malaysian Medical Association (channel + endorsement)
- **Technology** — Z.AI ILMU platform (production LLM), OpenAI (dev LLM + Whisper), Supabase (Postgres-as-a-service), Neo4j Aura (graph DB), Twilio + Meta Cloud (WhatsApp Business API), Sentry (observability)
- **Distribution** — Clinic-management VARs (medical-equipment distributors with existing clinic relationships), Malaysian pharmacy chains for cross-referrals

## 6.8 Risk Assessment

| Risk | Impact | Likelihood | Mitigation |
|---|---|---|---|
| **AI clinical inaccuracy** | High | Medium | Doctor-in-the-loop confirm step on every record; AI-vs-human visual distinction; full audit trail. |
| **Doctor adoption resistance** | High | Medium | 30-day no-commit pilot; on-site training; Hermes adapts to doctor style instead of forcing template change. |
| **Bahasa Malaysia STT accuracy** | Medium | Medium | Phase-1 text-first; Phase-2 voice with continuous fine-tuning on Malaysian clinic samples. |
| **PDPA compliance failure** | Critical | Low | Append-only audit log enforced at DB trigger; encryption at rest and in transit; annual third-party security audit; Malaysian data residency. |
| **LLM provider outage / price spike** | Medium | Medium | Pluggable OpenAI-compatible provider abstraction; multi-vendor failover (OpenAI ↔ Z.AI ILMU). |

# 7. Implementation & Financials

## 7.1 Validation Evidence (Current State)

| Validation type | Status as of 2026-05 |
|---|---|
| **Technical** | End-to-end MVP shipped: 4 portals (Patient / Doctor / Staff / Admin), 4 AI agents (Pre-Visit / Visit / Post-Visit / Evaluator), Postgres + Neo4j data layer, WhatsApp notifications via Twilio + Meta Cloud, Sentry across all 3 services, CI pipeline with security scans + Docker smoke tests. Evaluator soft-block verified at all 3 enforcement points. |
| **Clinical-safety** | Drug knowledge graph seeded with high-frequency Malaysian primary-care prescription set; six validators (DDI, allergy, dose, pregnancy, hallucination, completeness) live in the Evaluator agent; doctor override-with-reason audit trail enforced. |
| **Compliance** | PDPA audit log enforced at DB-trigger level (UPDATE/DELETE rejected); JWT-principal identity derivation on every per-patient mutation; cross-patient PHI leak post-mortem documented and remediated (`docs/post-mortem/2026-04-30-cross-patient-phi-leak.md`). |
| **Market** | 5–10 Klang Valley lighthouse clinics in pilot conversation; 3 University Malaya Faculty of Medicine partner clinics targeted for Phase-1 clinical validation; MMA channel-partnership discussion in progress. |
| **User** | Bahasa Malaysia + English code-switched intake demo flows verified; voice STT benchmarked at ~92% accuracy on Malaysian-accented English; staff walk-in registration and check-in flows verified end-to-end via Playwright. |

## 7.2 Go-to-Market Strategy

**City-by-city expansion.** Klang Valley (KL + Selangor) first — it concentrates ~40% of Malaysia's private GP clinics within a 60 km radius, supports on-site training economics, and offers proximity to MMA HQ for endorsement-driven growth.

**Three-stage adoption funnel:**

1. **Free 30-day pilot** with 5–10 lighthouse clinics, including 3 Faculty-of-Medicine partner clinics for clinical validation.
2. **Paid SME rollout** through MMA channel partnerships, targeted Facebook + LinkedIn campaigns to clinic owner-operators, and SEO content on documentation-burden topics.
3. **Chain expansion** in Year 2 — convert 3–5 multi-branch clinic chains to the Enterprise tier, bringing 60+ clinics each.

## 7.3 Implementation Plan

| Phase | Timeline | Action | Target |
|---|---|---|---|
| **Soft Launch** | Month 1–4 | Klang Valley pilots, lighthouse case studies, MVP hardening | 10 paying clinics, 90% doctor satisfaction in pilot |
| **Expansion** | Month 5–14 | Penang + Johor + Singapore launch, MMA channel activation, chain pilots | 250 paying clinics across 4 cities |
| **Optimization** | Month 15–24 | Hermes auto-learning loop, BM voice STT v2, multi-branch admin, ISO/IEC 27001 audit | 700 paying clinics, RM 9.5M ARR, cash-flow positive |
| **Total investment** | — | Seed RM 3.5M (engineering + GTM); Series A RM 10M for SEA expansion in Year 3 | Reach Year-3 ARR of RM 9.5M |

## 7.4 Unit Economics

| Metric | Value | Note |
|---|---|---|
| **CAC (blended)** | RM 800 / clinic | Field sales + MMA channel + digital, weighted by Y2 mix |
| **ARPC (Year 2 avg)** | RM 11,200 / year | Tier mix: 40% Starter, 50% Pro, 10% Enterprise |
| **Gross margin** | ~70% | LLM cost ~RM 0.40 / consultation; ~50 consultations / day / doctor |
| **Payback period** | ~3 months | RM 800 CAC / (RM 11,200 × 70% / 12) ≈ 3.0 months |
| **LTV (24-mo retention)** | ~RM 15,700 | RM 11,200 × 70% × 2 years |
| **LTV : CAC** | ~19.6× | Healthy SaaS benchmark is ≥3× |

## 7.5 Revenue Projections (3-Year)

| Metric | Year 1 | Year 2 | Year 3 |
|---|---|---|---|
| **Paying clinics (EOY)** | 50 | 250 | 700 |
| **Active cities** | 1 (Klang Valley) | 4 (KV, Penang, Johor, Singapore) | 6 (+ Ipoh, Kuching) |
| **Avg ACV (RM)** | 9,600 | 11,200 | 13,500 |
| **Subscription revenue** | RM 480K | RM 2.6M | RM 7.4M |
| **Metered AI revenue** | RM 90K | RM 480K | RM 1.5M |
| **Services revenue** | RM 30K | RM 120K | RM 600K |
| **Total revenue** | **RM 600K** | **RM 3.2M** | **RM 9.5M** |
| **OpEx** | RM 2.4M | RM 4.0M | RM 8.0M |
| **Net position** | (RM 1.8M) | (RM 0.8M) | **+RM 1.5M** |

## 7.6 Risk-Adjusted Scenarios

| Scenario | Y3 paying clinics | Y3 ARR | Trigger |
|---|---|---|---|
| **Bear** | 400 | RM 5.4M | MMA channel slips, slower BM-voice rollout, Singapore launch delayed to Y4 |
| **Base** | 700 | RM 9.5M | Plan as stated; cash-flow positive Month 28 |
| **Bull** | 1,000 | RM 13.5M | 2 chain enterprise wins (60+ clinics each), Singapore + Indonesia early entry |

## 7.7 Key Financial Milestones

- **Month 4** — 10 paying clinics in Klang Valley, RM 80K ARR, validated 50% doctor-time-saved benchmark
- **Month 12** — 50 paying clinics across 1 city, RM 600K ARR, MMA channel agreement signed
- **Month 18** — 250 paying clinics across 4 cities, RM 3.0M ARR, Series A close
- **Month 24** — 500 paying clinics, RM 6.0M ARR, gross-margin breakeven
- **Month 28** — Cash-flow positive
- **Month 36** — 700 paying clinics, RM 9.5M ARR, expansion review for Indonesia + Thailand